

# Bits, Bytes and Integers

Introduction to Computer Systems

<https://xjtu-ics.github.io/>

Danfeng Shan

Xi'an Jiaotong University

# Today: Bits, Bytes, and Integers

**Representing information as bits**

**Bit-level manipulations**

**Integers**

Representation: unsigned and signed

Conversion, casting

Expanding, truncating

Addition, multiplication, shifting

**Representations in memory, pointers, strings**

# Everything is bits

**Each bit is 0 or 1**

**By encoding/interpreting sets of bits in various ways**


Computers determine what to do (instructions)

... and represent and manipulate numbers, sets, strings, etc...

**Why bits? Electronic Implementation**

Easy to store with bistable elements (双稳态器件)

Reliably transmitted on noisy and inaccurate wires



# For example, can count in binary

## Base 2 Number Representation

Represent  $15213_{10}$  as  $11101101101101_2$

Represent  $1.20_{10}$  as  $1.0011001100110011[0011]..._2$

Represent  $1.5213 \times 10^4$  as  $1.1101101101101_2 \times 2^{13}$

# Encoding Byte Values

## Byte = 8 bits

Binary  $00000000_2$  to  $11111111_2$

Decimal:  $0_{10}$  to  $255_{10}$

Hexadecimal  $00_{16}$  to  $FF_{16}$

Base 16 number representation

Use characters '0' to '9' and 'A' to 'F'

Write  $FA1D37B_{16}$  in C as

- `0xFA1D37B`
- `0xfa1d37b`

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

15213: 0011 1011 0110 1101  
          3      B      6      D

# Example Data Representations

C Data Type	Typical 32-bit	Typical 64-bit
<code>char</code>	1	1
<code>short</code>	2	2
<code>int</code>	4	4
<code>long</code>	4	8
<code>float</code>	4	4
<code>double</code>	8	8
<code>pointer</code>	4	8

# Example Data Representations

C Data Type	Typical 32-bit	Typical 64-bit
<b>char</b>	1	1
<b>short</b>	2	2
<b>int</b>	4	4
<b>long</b>	4	8
<b>float</b>	4	4
<b>double</b>	8	8
<b>pointer</b>	4	8

# Example Data Representations

C Data Type	Typical 32-bit	Typical 64-bit
<b>char</b>	1	1
<b>short</b>	2	2
<b>int</b>	4	4
<b>long</b>	4	8
<b>float</b>	4	4
<b>double</b>	8	8
<b>pointer</b>	4	8

“ILP32”

“LP64”



# Today: Bits, Bytes, and Integers

Representing information as bits

**Bit-level manipulations**

**Integers**

Representation: unsigned and signed

Conversion, casting

Expanding, truncating

Addition, multiplication, shifting

**Representations in memory, pointers, strings**

# Boolean Algebra

## Developed by George Boole in 19th Century

Algebraic representation of logic

Encode "True" as 1 and "False" as 0

And

$A \& B = 1$  when both  $A=1$  and  $B=1$

$\&$	0	1
0	0	0
1	0	1

Or

$A | B = 1$  when either  $A=1$  or  $B=1$  or both

$ $	0	1
0	0	1
1	1	1

Not

$\sim A = 1$  when  $A=0$

$\sim$	0	1
	1	0

Exclusive-Or (Xor)

$A \wedge B = 1$  when  $A=1$  or  $B=1$ , but not both

$\wedge$	0	1
0	0	1
1	1	0

# General Boolean Algebras

## Operate on Bit Vectors

Operations applied bitwise

$\begin{array}{r} 01101001 \\ \& 01010101 \\ \hline 01000001 \end{array}$	$\begin{array}{r} 01101001 \\   01010101 \\ \hline 01111101 \end{array}$	$\begin{array}{r} 01101001 \\ ^ 01010101 \\ \hline 00111100 \end{array}$	$\begin{array}{r} \sim 01010101 \\ \hline 10101010 \end{array}$
---	--	--	---

# Bit-Level Operations in C

## Operations `&`, `|`, `~`, `^` Available in C

Apply to any “integral” data type

long, int, short, char, unsigned

View arguments as bit vectors

Arguments applied bit-wise

# Contrast: Logic Operations in C

## Contrast to Bit-Level Operators

Logic Operations: `&&`, `||`, `!`

View 0 as “False”

Anything nonzero as “True”

Always return 0 or 1

Early termination

## Examples (char data type)

`!0x41 → 0x00`

`!0x00 → 0x01`

`!!0x41 → 0x01`

`0x69 && 0x55 → 0x01`

`0x69 || 0x55 → 0x01`

`p && *p` (avoids null pointer access)

Watch out for `&&` vs. `&` (and `||` vs. `|`)...  
Super common C programming pitfall!

# Shift Operations

## Left Shift: $x \ll y$

Shift bit-vector  $x$  left  $y$  positions

- Throw away extra bits on left

Fill with 0's on right

## Right Shift: $x \gg y$

Shift bit-vector  $x$  right  $y$  positions

Throw away extra bits on right

Logical shift

Fill with 0's on left

Arithmetic shift

Replicate most significant bit on left

## Undefined Behavior

Shift amount  $< 0$  or  $\geq$  word size

Argument $x$	01100010
$\ll 3$	00010000
Log. $\gg 2$	00011000
Arith. $\gg 2$	00011000

Argument $x$	10100010
$\ll 3$	00010000
Log. $\gg 2$	00101000
Arith. $\gg 2$	11101000

# Today: Bits, Bytes, and Integers

Representing information as bits

Bit-level manipulations

## **Integers**

Representation: unsigned and signed, negation

Conversion, casting

Expanding, truncating

Addition, multiplication, shifting

**Representations in memory, pointers, strings**

# Question?

```
int foo = -1;  
unsigned bar = 1;  
  
(foo < bar) == true ?
```



# Encoding “Integers”

## Unsigned

Given a bit vector  $x$ ,  
 $w$  bits long...

$$\text{B2U}(x) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

## Signed (twos complement)

$$\text{B2T}(x) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

← **Sign Bit**

## Examples ( $w = 5$ )

$\pm 16$	8	4	2	1
0	1	0	1	0

$$0 + 8 + 0 + 2 + 0 = 10$$

16	8	4	2	1
1	0	1	1	0
-16	8	4	2	1

$$16 + 0 + 4 + 2 + 0 = 22$$

$$-16 + 0 + 4 + 2 + 0 = -10$$

# Negation: Complement & Increment

Negate through complement and increase

$$\sim x + 1 == -x$$

Why?

$$-x + x == 0 \text{ (by definition)}$$

$$\sim x + x == 1111\dots111 == -1$$

$$\sim x + x + 1 == 0$$

$$(\sim x + 1) + x == 0$$

$$\sim x + 1 == -x$$

$$\begin{array}{r} x \quad 10011101 \\ + \quad \sim x \quad 01100010 \\ \hline -1 \quad 11111111 \end{array}$$

Example:  $x = 15213$

	Decimal	Hex	Binary
<b>x</b>	<b>15213</b>	3B 6D	00111011 01101101
<b>~x</b>	<b>-15214</b>	C4 92	11000100 10010010
<b>~x+1</b>	<b>-15213</b>	C4 93	11000100 10010011
<b>y</b>	<b>-15213</b>	C4 93	11000100 10010011


# Complement & Increment Examples

$$x = 0$$

	Decimal	Hex	Binary
0	0	00 00	00000000 00000000
$\sim 0$	-1	FF FF	11111111 11111111
$\sim 0 + 1$	0	00 00	00000000 00000000

$$x = T_{\min}$$


	Decimal	Hex	Binary
$x$	-32768	80 00	10000000 00000000
$\sim x$	32767	7F FF	01111111 11111111
$\sim x + 1$	-32768	80 00	10000000 00000000



Oops!  
It's still  
negative!

Eight negative values:  
-1, -2, ..., -8

Mathematicians would prefer it if a 4-bit signed number could represent values -8...8, but that's  $2^4 + 1$  values, so they won't all fit.



Eight *non*-negative values:  
0, 1, ..., 7

What if we made a 4-bit signed number only represent values -7...7? Then we wouldn't be using bit pattern 1000...

# Today: Bits, Bytes, and Integers

Representing information as bits

Bit-level manipulations

Integers

Representation: unsigned and signed, negation

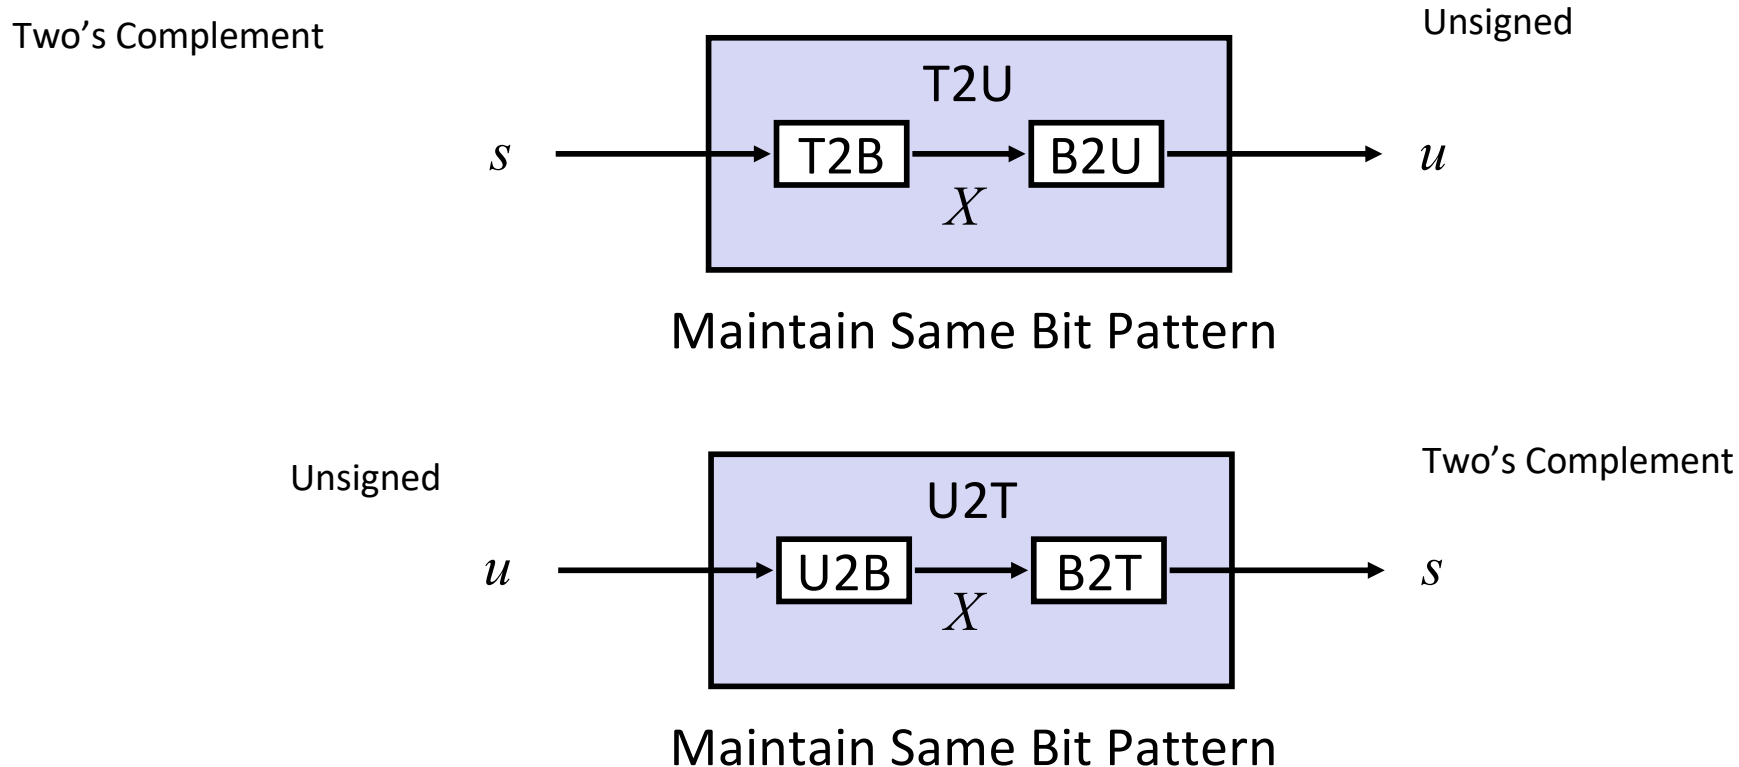
Conversion, casting

Expanding, truncating

Addition, multiplication, shifting

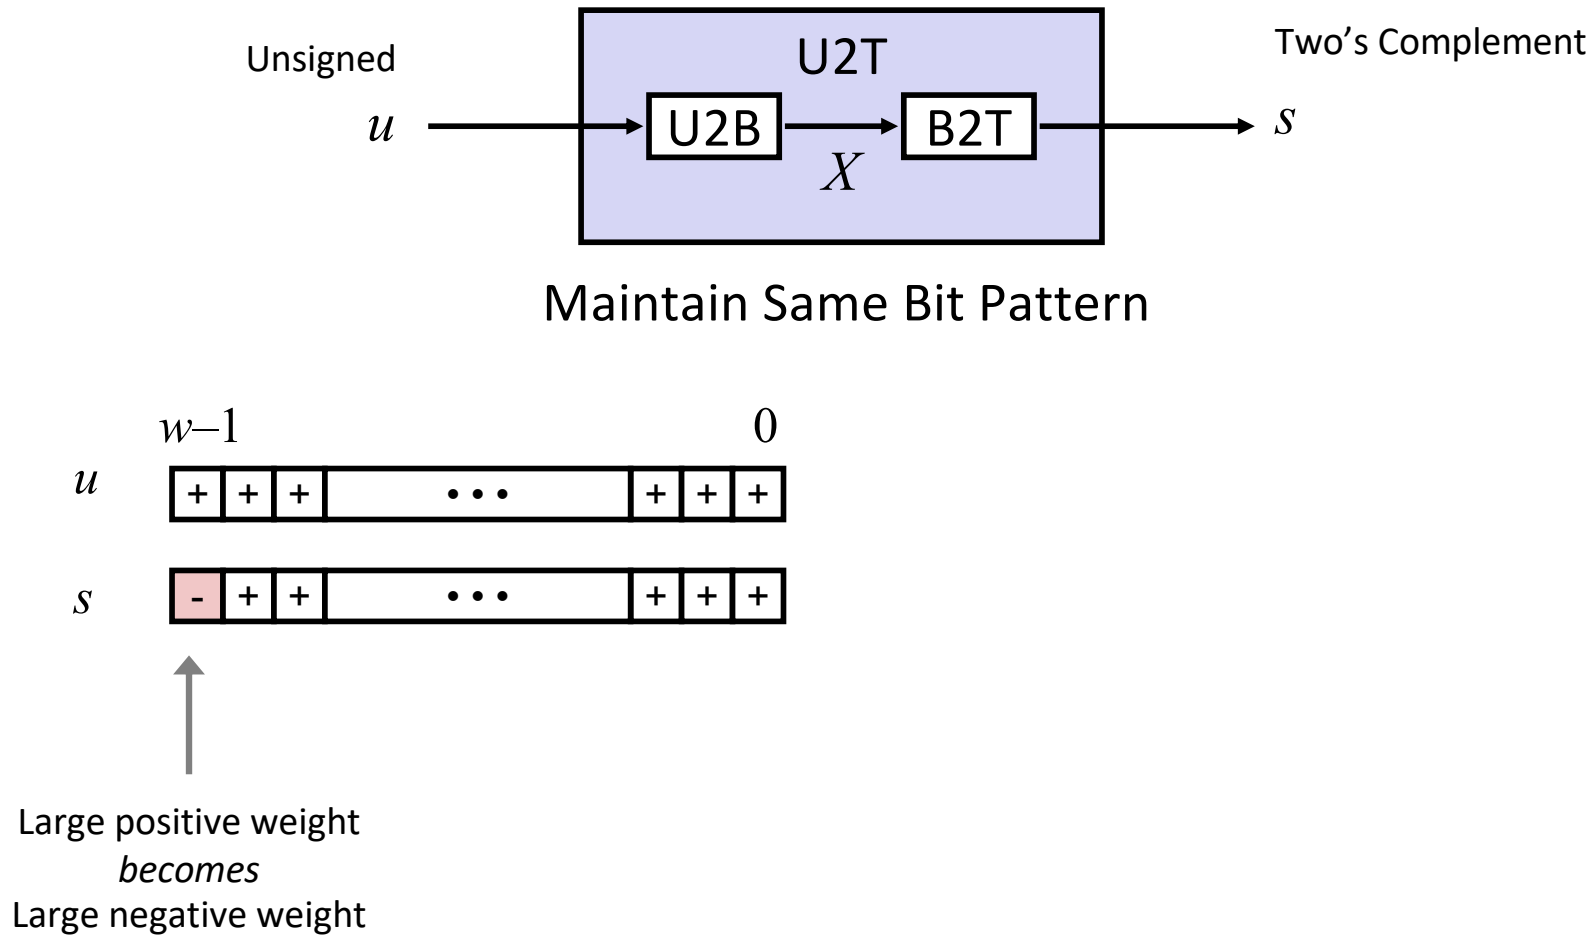
Representations in memory, pointers, strings

# Mapping Between Signed & Unsigned



Mappings between unsigned and two's complement numbers:  
**Keep bit representations and reinterpret**

# Relation between Signed & Unsigned



# Mapping Signed ↔ Unsigned

Bits	Signed		Unsigned
0000	0	$\longleftrightarrow$ =	0
0001	1		1
0010	2		2
0011	3		3
0100	4		4
0101	5		5
0110	6		6
0111	7		7
1000	-8	$\longleftrightarrow$ $\pm 16$	8
1001	-7		9
1010	-6		10
1011	-5		11
1100	-4		12
1101	-3		13
1110	-2		14
1111	-1		15




# Conversion Visualized

## 2's Comp. → Unsigned

Ordering Inversion

Negative → Big Positive

2's Complement  
Range



# Signed vs. Unsigned in C

## Constants

By default are considered to be signed integers

Unsigned if have “U” as suffix

`0U, 4294967259U`

## Casting

Explicit casting between signed & unsigned same as U2T and T2U

```
int tx, ty;  
unsigned ux, uy;  
tx = (int) ux;  
uy = (unsigned) ty;
```

Implicit casting also occurs via assignments and procedure calls

```
tx = ux;                int fun(unsigned u) ;  
uy = ty;                uy = fun(tx) ;
```

# Question?

example02.c

```
int foo = -1;  
unsigned bar = 1;  
  
foo < bar == true ?
```

# Casting Surprises

## Expression Evaluation

If there is a mix of unsigned and signed in single expression,  
*signed values implicitly cast to unsigned*

Including comparison operations `<`, `>`, `==`, `<=`, `>=`

Examples:

Constant 1	Constant 2	Relation	Evaluation
0	0U	==	Unsigned
-1	0	<	Signed
-1	0U	>	Unsigned
INT_MAX	INT_MIN	>	Signed
(unsigned) INT_MAX	INT_MIN	<	Unsigned
-1	-2	>	Signed
(unsigned) -1	-2	>	Unsigned
INT_MAX	((unsigned) INT_MAX) + 1	<	Unsigned
INT_MAX	(int) ((unsigned) INT_MAX) + 1	>	Signed

# Summary

## Casting Signed $\leftrightarrow$ Unsigned: Basic Rules

Bit pattern is maintained

But reinterpreted

Can have unexpected effects: adding or subtracting  $2^w$

**Expression containing signed and unsigned int**

`int` is cast to `unsigned`!!

# Today: Bits, Bytes, and Integers

Representing information as bits

Bit-level manipulations

**Integers**

Representation: unsigned and signed, negation

Conversion, casting

Expanding, truncating

Addition, multiplication, shifting

**Representations in memory, pointers, strings**

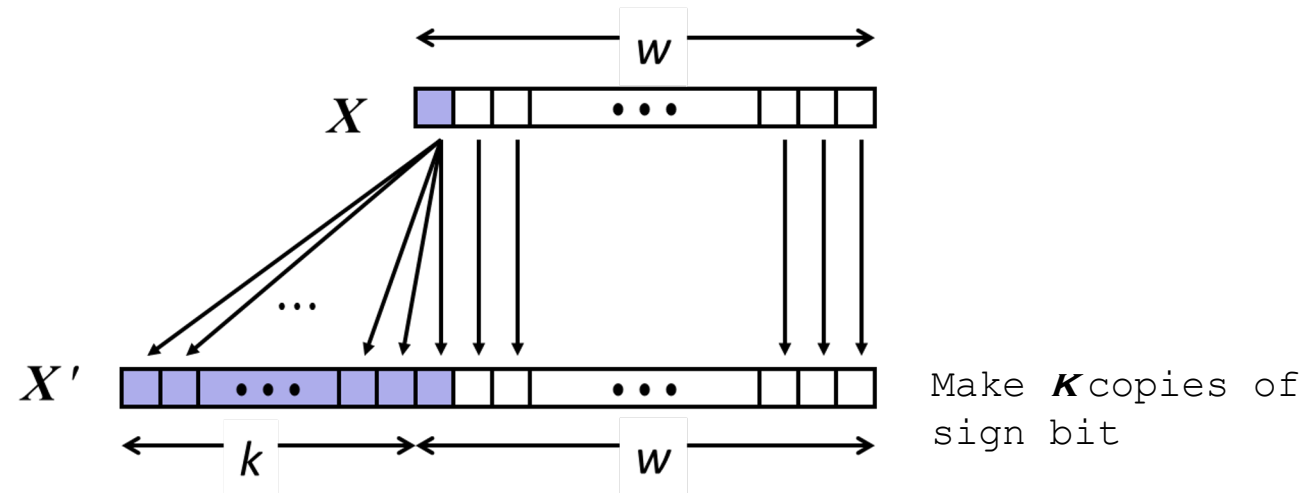
# Question?

example03.c

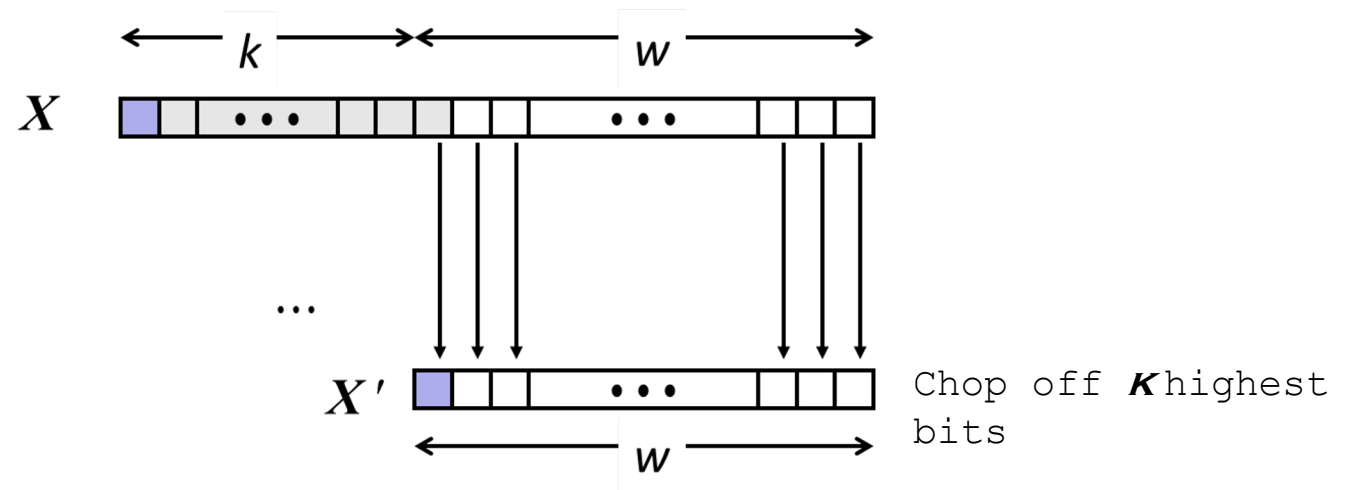
```
int x = 0x8000;  
short sx = (short) x;  
int y = sx;
```

# Sign Extension and Truncation

## Sign Extension



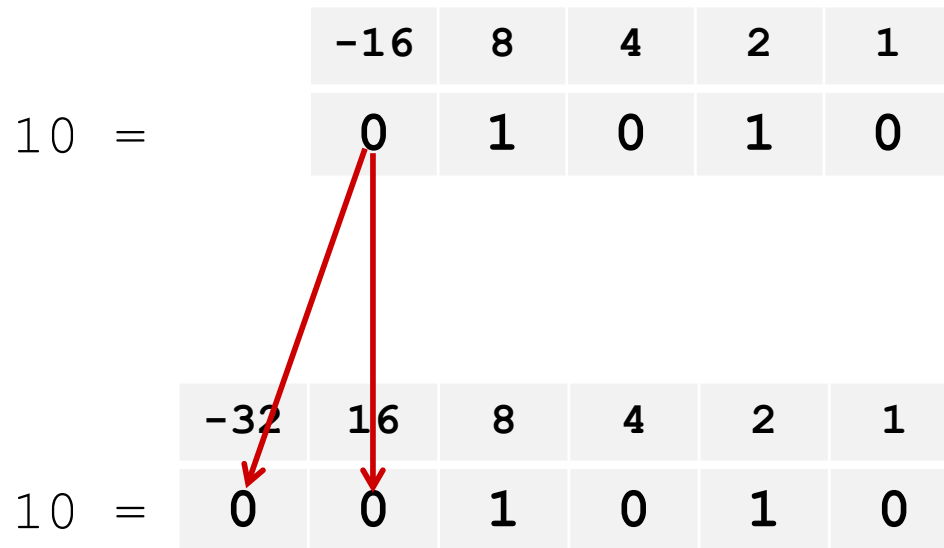
## Truncation



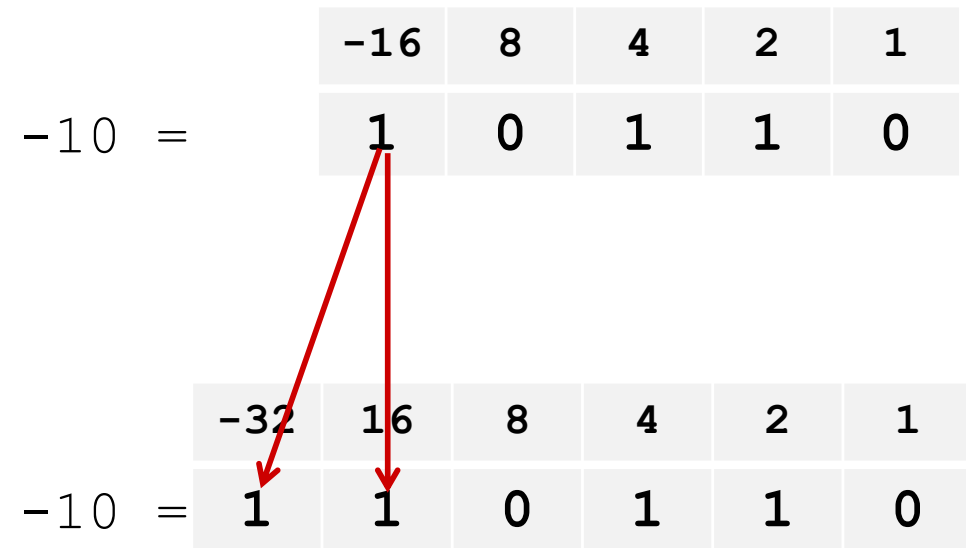


# Sign Extension: Simple Example

Positive number



Negative number



# Truncation: Simple Example

No sign change

	-16	8	4	2	1
2 =	0	0	0	1	0

	-8	4	2	1
2 =	0	0	1	0

	-16	8	4	2	1
-6 =	1	1	0	1	0

	-8	4	2	1
-6 =	1	0	1	0

Sign change

	-16	8	4	2	1
10 =	0	1	0	1	0

	-8	4	2	1
-6 =	1	0	1	0

	-16	8	4	2	1
-10 =	1	0	1	1	0

	-8	4	2	1
6 =	0	1	1	0

# Question?

example03.c

```
int x = 0x8000;  
short sx = (short) x;  
int y = sx;
```

# Today: Bits, Bytes, and Integers

Representing information as bits

Bit-level manipulations

Integers

- Representation: unsigned and signed

- Conversion, casting

- Expanding, truncating

- Addition, multiplication, shifting

Representations in memory, pointers, strings

# Unsigned Addition

Operands:  $w$  bits



True Sum:  $w+1$  bits



Discard Carry:  $w$  bits



## Standard Addition Function

Ignores carry output

Hex    Decimal    Binary

0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

unsigned char

1110 1001	E9	233
+ 1101 0101	+ D5	+ 213
<hr/>	<hr/>	<hr/>
<hr/>	<hr/>	<hr/>

# Unsigned Addition

Operands:  $w$  bits



True Sum:  $w+1$  bits



Discard Carry:  $w$  bits



## Standard Addition Function

Ignores carry output

Hex  
Decimal  
Binary

0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

unsigned char

```

      1110 1001
    + 1101 0101
    -----
  1 1011 1110
    -----
      1011 1110
  
```

```

      E9
    + D5
    -----
   1BE
    -----
      BE
  
```

```

      233
    + 213
    -----
   446
    -----
      190
  
```

# Visualizing (Mathematical) Integer Addition

## Integer Addition

4-bit integers  $u, v$


Compute true sum

$\text{Add}_4(u, v)$

Values increase

linearly with  $u$  and  $v$

Forms planar surface



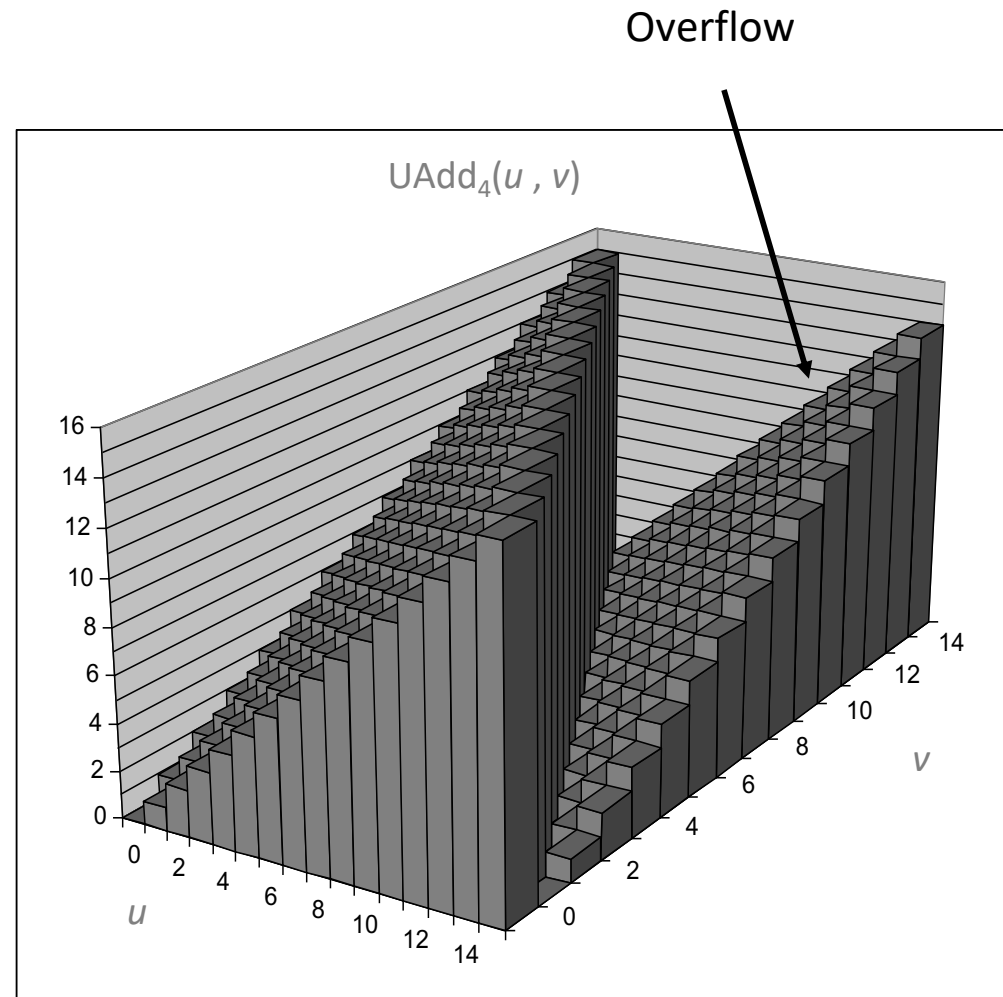
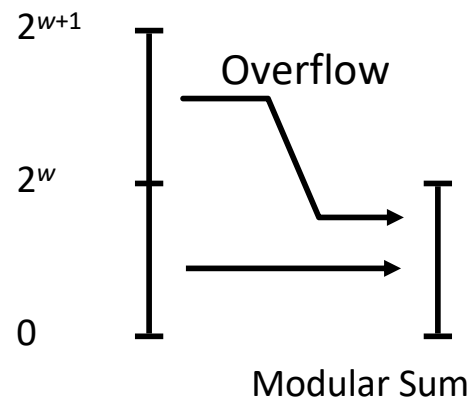
# Visualizing Unsigned Addition

## Wraps Around

If true sum  $\geq 2^w$

At most once

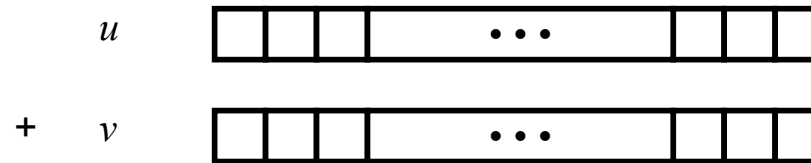
True Sum



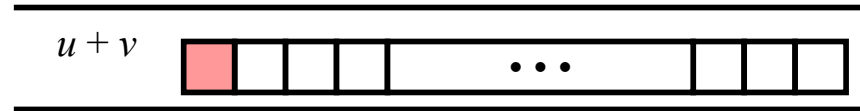


# Two's Complement Addition

Operands:  $w$  bits



True Sum:  $w+1$  bits



Discard Carry:  $w$  bits



## TAdd and UAdd have Identical Bit-Level Behavior

Signed vs. unsigned addition in C:

```
int s, t, u, v;
```

```
s = (int) ((unsigned) u + (unsigned) v);
```

```
t = u + v
```

Will give `s == t`

1110 1001	E9	-23
+ 1101 0101	+ D5	+ -43
<hr/>	<hr/>	<hr/>
1 1011 1110	1BE	-66
<hr/>	<hr/>	<hr/>
1011 1110	BE	-66

# Visualizing 2's Complement Addition

## Values

4-bit two's comp.

Range from -8 to +7

## Wraps Around

If  $\text{sum} \geq 2^{w-1}$

Becomes  
negative

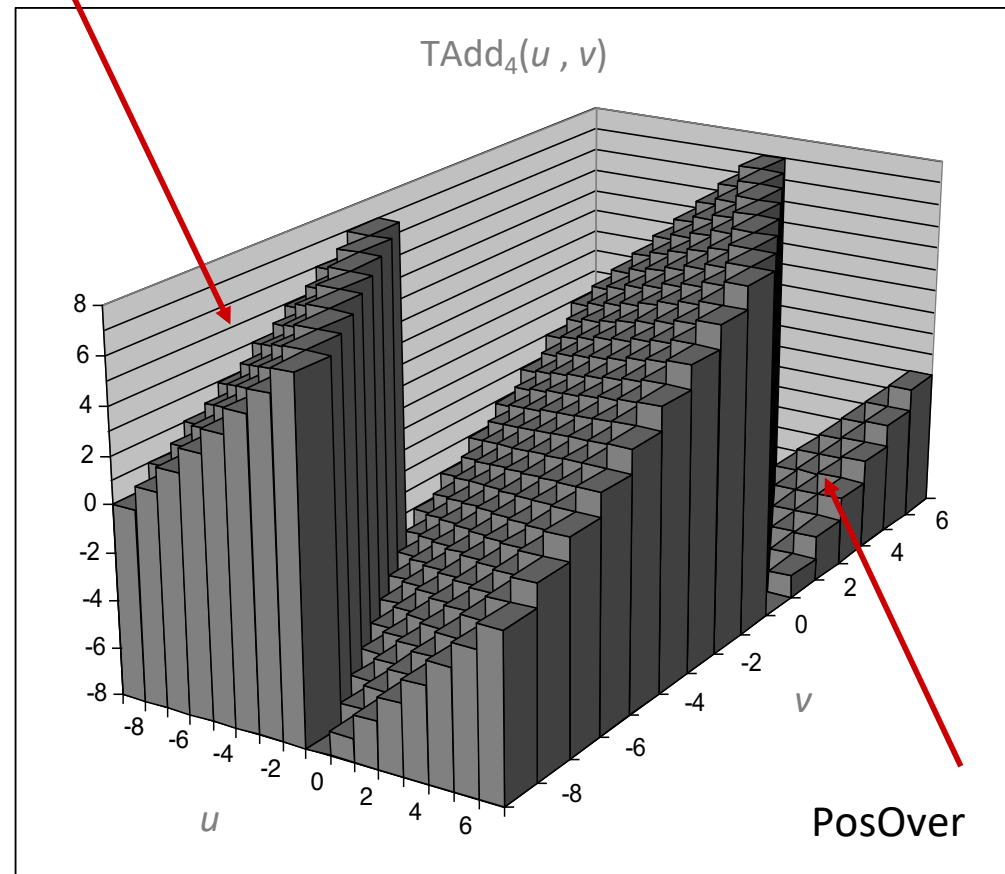
At most once

If  $\text{sum} < -2^{w-1}$

Becomes  
positive

At most once

NegOver




# TAdd Overflow

## Functionality

True sum requires  $w+1$  bits

Drop off MSB

Treat remaining bits as 2's comp. integer



# Today: Bits, Bytes, and Integers

Representing information as bits

Bit-level manipulations

**Integers**

Representation: unsigned and signed

Conversion, casting

Expanding, truncating

Addition, multiplication, shifting

**Representations in memory, pointers, strings**

# Shifting

## Left Shift: $x \ll y$

Shift bit-vector  $x$  left  $y$  positions

Throw away extra bits on left

Fill with 0's on right

Equivalent to multiplying by  $2^y$

## Right Shift: $x \gg y$

Shift bit-vector  $x$  right  $y$  positions

Throw away extra bits on right

Two kinds:

“Logical”: Fill with 0's on left

“Arithmetic”: Replicate most significant bit on left

*Almost* equivalent to dividing by  $2^y$

## Undefined Behavior (in C)

Shift amount  $< 0$  or  $\geq$  word size

Argument $x$	01100010
$\ll 3$	00010000
Logical $\gg 2$	00011000
Arithmetic $\gg 2$	00011000

Argument $x$	10100010
$\ll 3$	00010000
Logical $\gg 2$	00101000
Arithmetic $\gg 2$	11101000

# Multiplication

**Goal: Computing Product of  $w$ -bit numbers  $x, y$**

Either signed or unsigned

**But, exact results can be bigger than  $w$  bits**

Unsigned: up to  $2w$  bits

$$\text{Result range: } 0 \leq x * y \leq (2^w - 1)^2 = 2^{2w} - 2^{w+1} + 1$$

Two's complement min (negative): Up to  $2w-1$  bits

$$\text{Result range: } x * y \geq (-2^{w-1}) * (2^{w-1} - 1) = -2^{2w-2} + 2^{w-1}$$

Two's complement max (positive): Up to  $2w$  bits, but only for  $(TMin_w)^2$

$$\text{Result range: } x * y \leq (-2^{w-1})^2 = 2^{2w-2}$$

**So, maintaining exact results...**

would need to keep expanding word size with each product computed  
is done in software, if needed

# Unsigned Multiplication in C

Operands:  $w$  bits

$u$  

$*$   $v$  

True Product:  $2*w$  bits

$u \cdot v$  

$\text{UMult}_w(u, v)$

Discard  $w$  bits:  $w$  bits



## Standard Multiplication Function

Ignores high order  $w$  bits

```

      1110 1001
    * 1101 0101
    -----
1100 0001 1101 1101
    -----
      1101 1101
  
```

```

      E9
    * D5
    -----
    C1DD
      DD
  
```

```

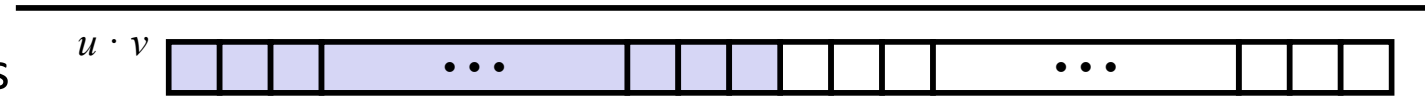
      233
    * 213
    -----
    49629
      221
  
```

# Signed Multiplication in C

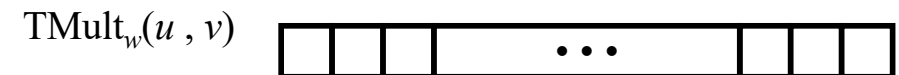
Operands:  $w$  bits



True Product:  $2*w$  bits



Discard  $w$  bits:  $w$  bits



## Standard Multiplication Function

Ignores high order  $w$  bits

Some of which are different for signed  
vs. unsigned multiplication

Lower bits are the same

$$\begin{array}{r}
 \phantom{*} \phantom{0000} \phantom{0011} 1110 \ 1001 \\
 * \phantom{0000} \phantom{0011} 1101 \ 0101 \\
 \hline
 0000 \ 0011 \ 1101 \ 1101 \\
 \hline
 \phantom{0000} \phantom{0011} 1101 \ 1101
 \end{array}$$

$$\begin{array}{r}
 \phantom{*} \phantom{03DD} \phantom{DD} \text{E9} \phantom{-23} \\
 * \phantom{03DD} \phantom{DD} \text{D5} \phantom{-43} \\
 \hline
 03DD \phantom{-989} \\
 \hline
 \phantom{03DD} \phantom{DD} \text{DD} \phantom{-35}
 \end{array}$$



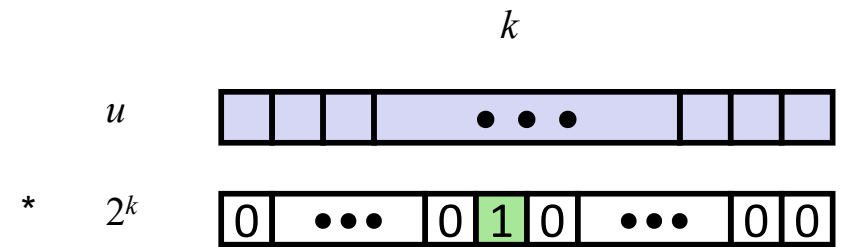
# Power-of-2 Multiply with Shift

## Operation

$u \ll k$  gives  $u * 2^k$

Both signed and unsigned

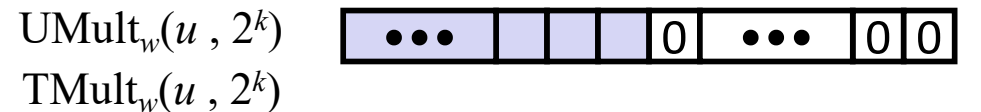
Operands:  $w$  bits



True Product:  $w+k$  bits



Discard  $k$  bits:  $w$  bits



## Examples

$u \ll 3 \quad == \quad u * 8$

$(u \ll 5) - (u \ll 3) \quad == \quad u * 24$

Most machines shift and add faster than multiply

Compiler generates this code automatically

# Today: Bits, Bytes, and Integers

Representing information as bits

Bit-level manipulations

**Integers**

Representation: unsigned and signed


Conversion, casting

Expanding, truncating

Addition, multiplication, shifting

**Representations in memory, pointers, strings**

# Byte-Oriented Memory Organization



## Programs refer to data by address

Imagine all of RAM as an enormous array of bytes

An address is an index into that array

A pointer variable stores an address

## System provides a private *address space* to each “process”

A process is an instance of a program, being executed

An address space is one of those enormous arrays of bytes

Each program can see only its own code and data within its enormous array

We'll come back to this later (“virtual memory” classes)

# Machine Words

## **Any given computer has a “Word Size”**

Nominal size of integer-valued data  
and of addresses

Historically, most machines used 32 bits (4 bytes) as word size  
Limits addresses to 4GB ( $2^{32}$  bytes)

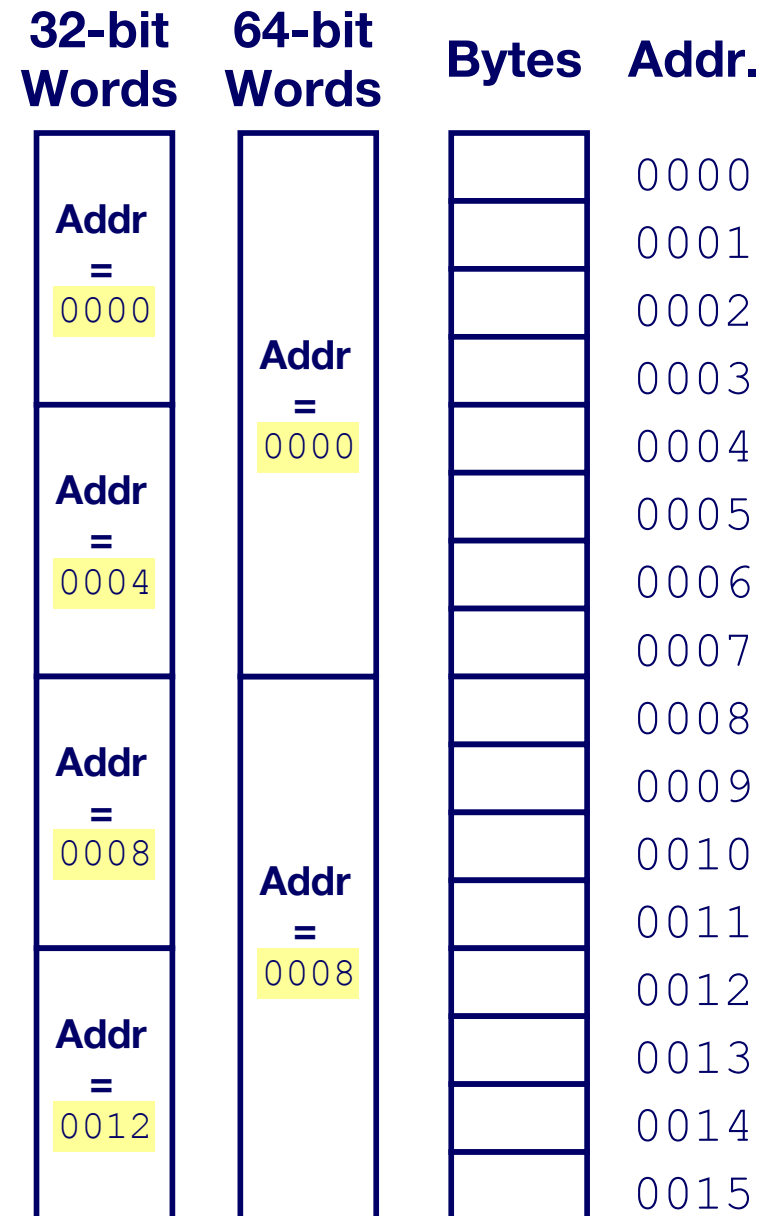
Currently, machines have 64-bit word size  
Potentially, could have 16 EB (exabytes) of addressable memory  
That's  $18.4 \times 10^{18}$  bytes

Machines still support multiple data formats  
Fractions or multiples of word size  
Always integral number of bytes

# Addresses *Always* Specify Byte Locations

Address of a word is address of the first byte in the word

Addresses of successive words differ by 4 (32-bit) or 8 (64-bit)



# Example Data Representations

C Data Type	Typical 32-bit	Typical 64-bit	x86-64
<code>char</code>	1	1	1
<code>short</code>	2	2	2
<code>int</code>	4	4	4
<code>long</code>	4	8	8
<code>float</code>	4	4	4
<code>double</code>	8	8	8
<code>pointer</code>	4	8	8

# Byte Ordering

**So, how are the bytes within a multi-byte word ordered in memory?**

## **Conventions**

Big Endian: Sun, PPC Mac, *network packet headers*

Least significant byte has highest address

Little Endian: *x86*, ARM processors running Android, iOS, and Windows

Least significant byte has lowest address


# Byte Ordering Example

## Example


Variable x has 4-byte value of 0x01234567

Address given by &x is 0x100

### Big Endian



### Little Endian






# Representing Integers

Decimal: 15213


Binary: 0011 1011 0110 1101

Hex: 3 B 6 D


`int A = 15213;`



`long int C = 15213;`



`int B = -15213;`



Two's complement  
representation

# Examining Data Representations

## Code to Print Byte Representation of Data

Casting pointer to unsigned char \* allows treatment as a byte array

```
typedef unsigned char *pointer;

void show_bytes(pointer start, size_t len){
    size_t i;
    for (i = 0; i < len; i++)
        printf("%p\t0x%.2x\n", start+i, start[i]);
    printf("\n");
}
```

### Printf directives:

%p:      Print pointer

%x:      Print Hexadecimal

# show\_bytes Execution Example

```
int a = 15213;  
printf("int a = 15213;\n");  
show_bytes((pointer) &a, sizeof(int));
```

## Result (Linux x86-64):

```
int a = 15213;  
0x7ffffb7f71dbc    6d  
0x7ffffb7f71dbd    3b  
0x7ffffb7f71dbe    00  
0x7ffffb7f71dbf    00
```

# Representing Pointers

```
int B = -15213;  
int *P = &B;
```

Sun	IA32	x86-64
EF	AC	3C
FF	28	1B
FB	F5	FE
2C	FF	82
		FD
		7F
		00
		00

**Different compilers & machines assign different locations to objects**

**Even get different results each time run program**

# Representing Strings

```
char S[6] = "18213";
```

## Strings in C

Represented by array of characters

Each character encoded in ASCII format

Standard 7-bit encoding of character set

Character "0" has code 0x30


– Digit  $i$  has code  $0x30+i$

String should be null-terminated

Final character = 0

## Compatibility

Byte ordering not an issue



# Representing x86 machine code

**x86 machine code is a sequence of *bytes***

Grouped into variable-length instructions, which look like strings...

But they contain embedded little-endian numbers...

## Example Fragment

Address	Instruction Code	Assembly Rendition
8048365:	5b	pop %ebx
8048366:	81 c3 <u>ab 12 00 00</u>	add \$0x12ab,%ebx
804836c:	83 bb 28 00 <u>00 00 00</u>	cmpl \$0x0,0x28(%ebx)

## Deciphering Numbers

Value:

Pad to 32 bits:

Split into bytes:

Reverse:

0x12ab

0x000012ab

00 00 12 ab

ab 12 00 00